

Section 2 evidence update — 25 September 2026

This is a follow-up to the 25 September leadership dossier, not a replacement for its dated observations. The comparison uses publicly described vendor functions, not independently tested vendor installations or a confirmed licence configuration. ORVIA evidence is from the customer-local codex-a00 synthetic profile. Equal standing is claimed only for the **bounded function** actually exercised, never for full platform breadth or production readiness.

What changed in ORVIA during this work window

Version 1 now has a non-model AI governance slice: tenant-scoped AI system and use inventory linked to purpose, processing activity, input asset, output system and processor; append-only risk, policy, control, evidence, monitoring, incident and separate-reviewer approval records; scoped API and workspace UI; reporting that distinguishes declaration from verification; and a durable input-inventory check job. The API integration run passed 32 assertions, including adverse and cross-tenant cases; the browser created and read back a synthetic use; and the actual monitor-worker run passed 9 assertions after the user authorized protected machine-identity setup. The monitor refuses to convert a staff-supplied OBSERVED label into independent evidence. No model or customer-data egress was added.

Function-level buying comparison

Function and public comparator signal	ORVIA evidence now	Position and next engineering proof
Discovery and data mapping. BigID markets automated discovery/classification; SecuriTI describes discovery across clouds and classification; OneTrust sells privacy data mapping. BigID , SecuriTI , OneTrust	Typed, provenance-aware scoped graph and allowlisted imports; linked AI use inventory.	The scoped mapping and provenance record is credible; autonomous estate discovery and broad source coverage are not equal. Next: allowlisted real read-only discovery connectors, observed provenance, freshness reconciliation and full source coverage tests.
Consent and notices. OneTrust describes consent/preferences and privacy operations; miniOrange and other India-oriented suites describe consent products. OneTrust , miniOrange	Exact notice-version consent, withdrawal epoch and replay protection, principal portal, immutable notice publication; existing synthetic tests.	A bounded customer-local grant and withdrawal journey exists. Touchpoint coverage, real deployment integrations, language/content legal review and sustained load evidence are needed for suite parity.
Rights, deletion and retention. OneTrust describes DSR fulfilment; BigID describes privacy and deletion functions. OneTrust , BigID	Rights state machine, hold checks, per-copy retention, safe allowlisted synthetic target actions and independent readback.	ORVIA's synthetic privacy-action journey has unusually explicit effect/unknown evidence, but no basis for broad-system parity. Next: real allowlisted connector, identity/search coverage, copy and processor readback, response package and adverse-case acceptance.
Vendors and assessments. OneTrust describes impact and vendor assessments; ServiceNow describes privacy assessments and processing activities. OneTrust , ServiceNow	Processor relationships, assessment findings and owner review; AI-use risk, policy, control and approval records.	The linked records provide a working subset. Next: complete DPIA templates/decision trail, third-party lifecycle, reassessment triggers and legal approval evidence.
Controls, evidence and continuous tests. Vanta, Drata, Sprinto and Scrut describe automated collection and monitoring. Vanta , Drata , Sprinto , Scrut	Durable workflow and action evidence, independent synthetic target readback, deterministic regression, and a tested durable AI inventory check that records a finding.	ORVIA can demonstrate a bounded independently observed privacy effect. It does not match those vendors' published integration or framework breadth. Next: real allowlisted integrations, periodic exact-candidate tests, delivered alerts and auditor export review.
Incidents and reporting. ServiceNow describes incident case workflows; OneTrust offers privacy operations reporting. ServiceNow , OneTrust	Incident time distinctions, scoped obligations and report subsets; AI incident links and finding counts.	A local traceable case/report subset exists. Next: reviewed regulatory content, real notification transport, delivery readback, report reconciliation and release acceptance.

Function and public comparator signal	ORVIA evidence now	Position and next engineering proof
Regulatory and GRC catalogs. MetricStream describes automated regulatory change and obligation mapping; ServiceNow documents privacy risk/control libraries. MetricStream , ServiceNow	Fixed policy and control records, deterministic test and evidence chain.	There is no equivalent maintained regulatory content or enterprise control library. Next: licensed/customer-reviewed sources, versioned applicability decisions, control mappings, update provenance and human legal approval.
Non-model AI governance. OneTrust and Securiti describe AI inventory, risk, monitoring and wider discovery/controls. OneTrust , Securiti	A working declared inventory, links, assessments, review, event history, incident and report API/UI; durable worker ran and recorded a finding.	Equal standing is limited to basic governed recording and human approval , subject to the synthetic test limit. Discovery of undeclared AI, model/runtime evaluation and input/output control enforcement are outstanding. ORVIA's absence of a model dependency is a deliberate customer-local product boundary, not a claim that competing AI functions are unnecessary.
Customer-local and commercial operations. Published SaaS features and packages vary by quote and deployment.	Local operational boundary, synthetic onboarding/licence/support/update subsets.	Do not infer lower price, certification or production qualification. Next: customer-held recovery decision, real packaging/install rehearsal, support boundary, vendor quotes and total-cost comparison.

Platform reading of Section 2

ORVIA now has tangible local privacy-action and non-model AI governance journeys, including a narrow, independently read-back synthetic privacy effect. The published discovery, integration, framework and enterprise GRC breadth of OneTrust, BigID, Securiti, Vanta, Drata, Sprinto, Scrut, ServiceNow and MetricStream remains broader than ORVIA's verified scope. ComplyIQ publicly describes data mapping, ROPA, vendors and approvals; the other India-oriented and specialist names still require module- and deployment-specific demonstrations. ORVIA cannot responsibly claim platform superiority or lower cost from index listings. [ComplyIQ](#)

The highest-value next proof is one approved real, read-only discovery integration followed by one safe controllable target action, independent readback, durable reconciliation, alert, evidence export and browser review on an identified candidate. The DPDP extension lane owns its separate paths; any shared application change requires coordination. The full T01–T34 application acceptance has **not** been produced for this work tree.